

PSDF: Particle Stream Data Format for N-Body Simulations

Will M. Farr^{a,*}, Jeff Ames¹, Piet Hut^c, Junichiro Makino^b, Steve McMillan¹, Takayuki Muranushi¹, Koichi Nakamura¹,
Keigo Nitadori¹, Simon Portegies Zwart¹

^a*Northwestern University Center for Interdisciplinary Research in Astrophysics, 2145 Sheridan Rd., Evanston IL 60208 USA*

^b*Interactive Research Center of Science, Graduate School of Science and Engineering Tokyo Institute of Technology, 2-12-1 Ookayama, Meguro, Tokyo 152-8551, Japan*

^c*Institute for Advanced Study, Princeton, NJ 08540, USA*

Abstract

We present a data format for the output of general N-body simulations, allowing the presence of individual time steps. By specifying a standard, different N-body integrators and different visualization and analysis programs can all share the simulation data, independent of the type of programs used to produce the data. Our Particle Stream Data Format, PSDF, is specified in YAML, based on the same approach as XML but with a simpler syntax. Together with a specification of PSDF, we provide background and motivation, as well as specific examples in a variety of computer languages. We also offer a web site from which these examples can be retrieved, in order to make it easy to augment existing codes in order to give them the option to produce PSDF output.

Keywords: Stellar dynamics, Method: N-body simulation

1. Introduction

The simplest N-body calculations use a shared time step length for all particles. This implies a rather simple structure of the output. With N particles and k time steps, the output takes on the form of an $N \times k$ matrix of particle data, where the latter typically contain the mass, position and velocity of a single particle at a specific time, with possible additional information such as higher derivatives of the position (acceleration, jerk, etc.), the value of the potential at the position of that particle, and so on. The output of this matrix can be done by ordering in time or by ordering by the identity of particles, in which case each world line is output separately.

Some complications may occur when particles are removed, for example because they are escaping from the system, or because they represent a star that undergoes a destructive supernova leaving no remnant. However, the basic I/O structure is simple enough that it is easy to present these kinds of data in one of the standard data formats, such as FITS or HDF, with a brief description of what is what.

The situation gets vastly more complicated, though, when we allow for individual time steps. Simulations of dense stellar systems, such as open and globular star clusters, as well as galactic nuclei, have relied on the use of individual time steps very early on, at least since the 1960s. The reason is that the presence of close binaries and triples in such systems would increase the computer power needed

by orders of magnitude in case of shared time steps, compared to individual time steps. In addition, cosmological codes, too, are moving toward the use of individual timesteps, given the increasingly large discrepancies of intrinsic time scales that come with increasingly high spatial resolution.

The simplest way to output data from individual time step codes would be to stick to shared time steps. Indeed, typical legacy codes, such as NBODY6, do just that by default. If all one wants to do is to make a fixed movie of a simulation run, that approach suffices. However, when we interactively inspect the results of a simulation run, we want to be able to zoom in and out, and speed up and slow down the rate at which we run the graphics presentation of the run. With a fixed initial output rate, it may not be possible to interpolate the motion of the particles that move at high speeds. Phrased differently, an output rate high enough to faithfully present the motion of all particles may be prohibitively expensive in terms of memory. It would be much better to let the graphics program itself decide how and where to extrapolate, given the original data it has received from a simulations code.

For example, when we display the dense center of a star cluster, the graphics program can then use the full information for the rapidly moving particles, while interpolating the data for the slower halo particles. Such an approach can easily save orders of magnitude of memory storage requirement. An implementation of this approach was made by Steve McMillan (199? – **Steve, please provide reference and a few-line summary**). However, this implementation was handcrafted for a specific code,

*Corresponding author

Email address: w-farr@northwestern.edu (Will M. Farr)

reflecting the data structure used in that code. Clearly, it would be desirable to have a more universal data format that would allow different codes to share data in a more transparent way.

Other concerns are to make a data format standard machine independent, to make allowance for parallel processing, and to avoid serious overhead penalties with respect to performance (**Jun, do you want to add a few lines here?**)

In the following sections, we describe a machine-independent algorithm-independent data format for storing the results of a simulation of point-mass gravitational dynamics using individual timesteps, which we call the “Particle Stream Data Format,” or PSDF.

2. Basic idea

We wish to store the evolving state of a gravitating system of N bodies throughout a simulation with individual timesteps for each particle. Conceptually, what we need is a stream of phase-space information of particles, updated each time the integration algorithm adjusts a particle’s phase space information. One possibility for such a stream could be:

```
particle_id, time, mass, x, y, z, vx, vy, vz, ...
particle_id, time, mass, x, y, z, vx, vy, vz, ...
particle_id, time, mass, x, y, z, vx, vy, vz, ...
```

However, the data format should be flexible enough to be able to include more information, if available, such as

- radius, and other info related to stellar evolution
- merger history
- fluid properties if a particle is an SPH particle
- **other things?**

One way to construct such a flexible data format is to use self-describing data format, such as XML or YAML. For simplicity, we adopt YAML (YAML, 2011) here; there are libraries for reading and writing YAML in many popular programming languages, and the format is simple enough to be understood easily by humans, even if they are not already familiar with it.

2.1. Some basics of YAML

The following is a simple example of data in YAML format.

```
--- !Particle
id: 0
r:
  - 0.1
  - 0.2
  - 0.3
v:
```

```
- -1
- -2
- -3
m: 1.0
```

In the above example, the line

```
--- !Particle
```

Is the header, which indicates that it describes the data of an object of type `Particle`. The line

```
id: 0
```

defines a field with name “id”, and value 0. The text

```
r:
  - 0.1
  - 0.2
  - 0.3
```

means the field “r” is an array with three elements. The first “-” means this line is a data for an array. By default, numbers without “.” are regarded as integers, and with “.” floating point. Note that indentation has meaning here and “-” must be indented the same level or deeper than “r” and should be aligned. The “v” and “m” fields behave similarly. The order of the fields is not important; the `Particle` object behaves as a *map* from names, like “r”, to values, like the array “[0.1, 0.2, 0.3]”.

3. Particle Stream Data Format

With the minimal description of YAML in the previous section, we can now define the Particle Stream Data Format: the data format is a stream of YAML representations of particle objects. For example, a valid PSDF fragment is

```
--- !Particle
id: 0
t: 0
r:
  - 0.1
  - 0.2
  - 0.3
v:
  - -1
  - -2
  - -3
m: 1.0
--- !Particle
id: 1
t: 0
r:
  - 0.2
  - 0.3
  - 0.4
v:
```

Name	Meaning
id	index (can be arbitrary text)
m	mass
t	time
t_max	max time to which this record is valid
r	position, array with three elements
v	velocity, array with three elements
pot	potential
acc	acceleration, array with three elements
jerk	jerk, array with three elements
snap	snap, array with three elements
crackle	crackle, array with three elements
pop	pop, array with three elements

Table 1: The reserved names of the PSDF and their meanings. Other names occurring in the data stream are to be ignored if they are not meaningful to an application or interpreted in an application-specific way if they are. A particle object in a PSDF data stream need not contain any of these names, but if it does they must have the meanings above; similarly, if a particle object contains data with the above meanings, it should be identified by the associated standard name.

```

- 0
- 0
- 0
m: 1.0

```

This fragment describes two particles, at $t = 0$, with ids 0 and 1.

Particle objects behave in YAML as mappings from names to values. Therefore, a specification of the meaning of certain names and a procedure for handling unknown names in the stream are sufficient to define the data format. In Table 1 we list the reserved names of our PSDF. Any particular particle in a PSDF stream need not include a value in its map for any of these names, but if it does, the value must have the meaning in the table; similarly, if a particle object in a PSDF stream does contain a value with one of the meanings in the table, then it should be identified by the corresponding name. Note that these requirements allow for easy extension of a PSDF stream with application-specific information by including any names and values not in Table 1 needed by the specialized application. Programs that understand this additional information can benefit, while those that do not will still be able to function using the basic information from any included values of Table 1. In the future, we intend to provide extensions that are useful for hierarchical decomposition of an N -body system and for description of fluid SPH particles.

We require that time, position, velocity and higher derivatives are consistent (for example, if position is given in parsecs and time in years, velocity must be in parsec/year). The name `t_max` is rather special, in that it gives the maximum possible time that this record is used to predict the orbit of this particle; we expect that this may prove useful to prevent invalid extrapolation in pro-

grams that process the PSDF.

A complete PSDF object is a stream of particle objects describing the states of individual particles in the system at particular times. Such a stream may be consumed as it is produced, as in the case of an integrator program whose output is directed to a graphics program that displays the result of the integration; or such a stream may be written to one or more files to be processed at a later date. We do not impose any particular ordering on the particle records in a PSDF stream. For some applications an ordering in time may be appropriate, while for others an ordering in particle id may be better, or even more complex orderings; later in this paper we provide references to code that can convert between the time and particle-id orderings.

4. Rationale

Our goal is to describe a data format that is

1. Space-efficient for storing the data individual-timestep N -body simulations.
2. Simple enough to be human readable and writeable.
3. Information-rich. It should allow for post-processing and analysis or even continuation or re-running of a simulation.
4. Flexible enough to accommodate the special needs of programs that have more complex objects than point-mass particles.
5. Simple and safe to read and parse within a program.
6. Composable, so that fragments of the data can be split off for separate analysis and recombined easily.

In this section we describe how these design goals led to the specification in the last section.

4.1. Space Efficient and Human-Readable

As outlined in the introduction, it is extremely wasteful to produce a complete system snapshot in an N -body simulation every time some of the bodies update their positions or velocities. The PSDF format allows for the output of only the changed data—the new states of the updated particles at the new time—to the stream. Though the native format is text-based, for readability, we have found that common compression algorithms such as `gzip` applied to files containing PSDF data produce output that is within 10% of the size of equivalent compressed binary data.

4.2. Information-Rich and Flexible

The fundamental state of a point-particle can be specified in 8 numbers: one mass, one time, three position coordinates and three velocity coordinates. However, some auxiliary information about the particle’s state can be very helpful: accurate prediction of the particle’s position and velocity—for example, to display its track in a visualization—can be facilitated by information about the higher derivatives of its position in time. When available

to the integration routine, these can be easily provided by our format (see Table 1).

Not all particles in many N -body simulations are point-masses! For example, simulations may attempt to model stellar evolution, and therefore store “star” properties like radii and masses, or entropy profiles with their particles. Or, simulations may include fluid particles subject to non-gravitational forces for SPH calculations. To attempt to standardize names for every possible particle property would result in a rigid and cumbersome format; instead, by allowing arbitrary application-specific names in PSDF particle mappings that can be ignored when not understood we permit complex applications to work with application-specific data while ensuring that simple applications can make use of the parts of the data they understand.

4.3. Simple, Safe, and Composable

By using a general format (YAML) in wide use for our PSDF, we ensure that there are mature, debugged libraries available for reading and writing our format. However, YAML is simple enough that it can be written and modified easily “by hand” in a text editor. The idea of streaming updates to individual particle states also meshes nicely with the evolution algorithms in most N -body integrators, making the format easy to write from within such a code.

PSDF objects are descriptions of *data*, not instructions for actions for an application. For example, there is no instruction for “adding” or “deleting” a particle from the stream. Adding such instructions would require applications to implement interpreters for implementing the instructions in the stream, which raises issues of security and language design that would significantly complicate the specification.

Finally, PSDF documents are *composable*, meaning that any two PSDF streams can be concatenated to form another valid PSDF stream, and a single stream can be split into any number of valid sub-streams. It is quite simple to write a short script in any number of languages that consumes a PSDF stream and produces another recording the history of a particular particle in the original stream, for example. Adding header information, or instructions, or any other meta-data to the stream defeats this goal by requiring specification of some way to split and combine the associated meta-data.

5. Repository

We have examples of codes that generate and manipulate PSDF streams at

<https://github.com/jmakino/Particle-Stream-Data-Format>

The examples are in various different languages, and include

- Code and references to several different individual-time-step integrators that can take PSDF input and advance the corresponding system in time, producing PSDF output at each step.
- Various post-processing tools that compute useful system properties from PSDF input.
- A visualization program that takes PSDF input and produces a 3D representation of the system that can be played forward and backward, zoomed, etc.

We hope that the examples we provide will make it easy for the community to use PSDF in their simulations, and that these users will, in turn, contribute their useful programs back to the repository as examples for future users.

Acknowledgments

YAML, 2011. The official yaml web site. [Http://www.yaml.org/](http://www.yaml.org/).